

Conditional GANs

Pokemon's type generation from scratch - Deep Learning course's exam

Matteo Robbiati

July 2022

GANs

➡ Generative Adversal Networks (aka GANs) are models composed of two neural networks:

- a **generator** G , which goal is to understand the data distribution ρ_{ds} . In order to do that, it takes random vector z from a latent space \mathcal{Z} ¹ with a probability density function ρ_z ² and maps it into an element with the same dimensionality of the data sample, building a fake sample space $\hat{\Omega}$, e.g. images in the form (64, 64, 3):

$$G : z \in \mathcal{Z} \rightarrow G(z|\vec{\theta}_g) \in \hat{\Omega}.$$

- a **discriminator**, D , which goal is to classify an input variable x as *generated* (0) or *real* (1):

$$D : x \in \{(64, 64, 3)\} \rightarrow D(x|\vec{\theta}_d) \in \{0, 1\}$$

¹with a lower dimensionality than the sample space

²typically in a simple way, e.g. through an uniform or a normal distribution

The loss function to optimize

➡ G wants D classifies all its purposes as 1, instead D wants to be able to classify correctly every input it receives. To summing up this conflict between the two network we build a **min-max game**:

$$\min_{\vec{\theta}_g} \max_{\vec{\theta}_d} \left\{ \langle \log D(x|\vec{\theta}_d) \rangle_{x \sim \rho_{ds}} + \langle [1 - \log D(G(z|\vec{\theta}_g)|\vec{\theta}_d)] \rangle_{z \sim \rho_z} \right\} \quad (1)$$

```
1 # the pseudo-algorithm
2 for n epochs do:
3     for k steps do:
4         let D classify an hybrid batch of data
5         update the discriminator weights with a B-P
6         generate a new fake batch using G
7         freeze D weights and classify the elements
8         update the generator weights with a B-P
```

➡ This idea of optimization is applied also in the case of conditional generative adversal nets.

Conditional GANs

A conditioned problem

- ➡ Conditional GANs were born for implementing a model that was able to generate conditioned data.
- ➡ The idea is to feed both G and D with an extra information y , which can be arbitrarily implemented³.

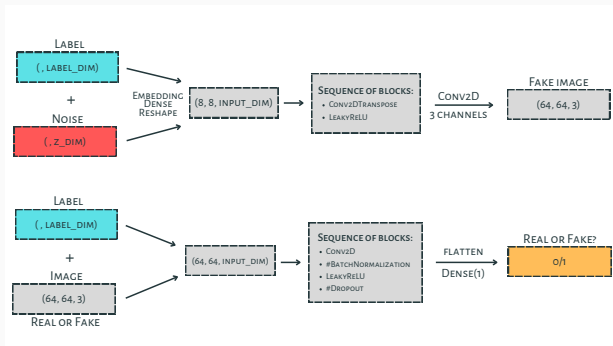


Figure 1: Above, G diagram. Below, D diagram.

³I treated two methods. In method 🖐 i used functional API mode for passing G and D two inputs, with method 🖐 i merged all inputs into a single array and built sequential models

This problem

- ➡ I want to generate 64×64 **RGB images** of pokemon of three different types: *grass*, *water* and *fire*.
- ➡ i managed [this dataset](#), which contained initially $\sim 17.3k$ images:
 - **deleted** all folders containing pokemon **different** from *water*, *grass* or *fire* (from 149 to 50 directories);
 - **standardized** the images to a 64×64 RGB and **png* format (removed 208 images incompatible with the RGB-png conversion);
 - **removed** all the images **without a white background**. For this step I've written a function which evaluate the percent value of *whiteness* of an image (removed 2450 images);
- ➡ At the end of the pre-processing i found a dataset of 1064 images divided into three type-classes, which i put into a `tf.keras.Dataset`.

A random sampling



How i built the dataset

➡ Since `cgan_ds` contains only a few elements, i have decided to apply a data augmentation strategy during the process of `ds` definition:

```
1 #a bit of data augmentation
2 data_augmentation = tf.keras.models.Sequential(
3     [
4         tf.keras.layers.RandomFlip("horizontal"),
5         tf.keras.layers.RandomRotation(0.1),
6         tf.keras.layers.RandomZoom(0.1),
7     ]
8 )
9
10 #prepare function
11 def prepare(ds, shuffle=True, augment=True):
12     ds.cache()
13     if shuffle:
14         ds = ds.shuffle(6000)
15     ds = ds.batch(batch_size)
16     if augment:
17         ds = ds.map(lambda x, y: (data_augmentation(x), y), num_parallel_calls=BUFF_SIZE)
18     return ds.prefetch(buffer_size=BUFF_SIZE)
19
20 #prepare my dataset as decided
21 dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))
22 dataset = prepare(dataset)
```

A random sampling



Conditioning the generator - Method

With the first way we want to embed all infos into **sheets**, from which we will get a (64,64,3) image.

1. the label is passed as integer and embedded into a vector with an arbitrary dimension which we call **embed_label**. Then, using a Dense layer composed of 8×8 nodes and a Reshape layer we transform the **embed_label** into a (8,8,1) sheet;
2. **noise** z is passed as a random vector of dimension 100, from which we build N_s sheets using a Dense with $N_s \times 8 \times 8$ nodes and a Reshape layer;
3. with a Concatenate layer we **combine the sheets** composing all the amount of input information, in the form $(8, 8, N_s + 1)$.

```
1 #embedding the label
2 input_label = Input(shape=(1,))
3 embed_label = Embedding(num_classes, 100)(input_label)
4 nodes = 8 * 8
5 embed_label = Dense(nodes)(embed_label)
6 embed_label = Reshape(8, 8, 1)(embed_label)
7 #embedding the noise
8 latent_input = Input(shape=(latent_dim,))
9 nodes = 64 * 8 * 8
10 gen = Dense(nodes)(latent_input)
11 gen = Reshape((8, 8, 64))(gen)
```

Conditioning the generator - Method 🙌

We encode all the information (both noise and class label) into a single input vector, from which we get the same sheets.

1. Once calculated the **input dimension**: $N_{in,G} = N_{classes} + z_{dim}$;
2. we transform the input information using a Dense layer with $8 \times 8 \times N_{in,G}$ nodes and reshape the output into sheets: $(8, 8, N_{in,G})$.

```
1 generator = keras.Sequential(  
2     [  
3         #embedding the input information  
4         #generator_in_channels is our N_in,G  
5         InputLayer((generator_in_channels,)),  
6         Dense(8 * 8 * generator_in_channels),  
7         LeakyReLU(alpha=0.2),  
8         Reshape((8, 8, generator_in_channels)),  
9         #(...)  
10    ]  
11 )
```

We operate in terms of sheets again:

1. the **label** is passed as first input and embedded in the same way of *G*, with the difference that it is transformed into a sheet of the form (64, 64, 1);
2. the **image** is provided as second input and pasted with the *label sheet* through a *Concatenate* layer.

```
1 # embedding the label
2 input_label = Input(shape=(1,))
3 embed_label = Embedding(1, 100)(input_label)
4 nodes = 64 * 64
5 embed_label = Dense(nodes)(embed_label)
6 embed_label = Reshape((64, 64, 1))(embed_label)
7 # passing the second input (the RGB image)
8 input_image = Input(shape=input.shape)
9 # concatenating the two inputs
10 merged_image = Concatenate()([input_image, embed_label])
```

We build up a sequence of sheets in the form (64, 64,) containing the image and the label information.

1. the **label** compose three sheets for a total dimension (64, 64, 3);
2. the **image** compose other three sheets in the same form (64, 64, 3);
3. all inputs are concatenated into six sheets as (64, 64, 6).

```
1 discriminator = keras.Sequential(  
2     [  
3         #discriminator_in_channels is (n_channels + n_classes)  
4         #so it is an input (64, 64, 6) in this case  
5         keras.layers.InputLayer((64, 64, discriminator_in_channels)),  
6         # Convolutions (...)  
7     ]  
8 )
```

- ➡ Conv2D and Conv2DTranspose operations respectively for reducing and increasing the data size

```
1 #discriminator's core
2 layers.Conv2D(512, (4, 4), strides=(1, 1), padding="same"),
3 layers.LeakyReLU(alpha=0.2),
4 layers.Conv2D(256, (4, 4), strides=(1, 1), padding="same"),
5 layers.LeakyReLU(alpha=0.2),
6 layers.Conv2D(128, (4, 4), strides=(1, 1), padding="same"),
7 layers.LeakyReLU(alpha=0.2),
8 layers.Conv2D(64, (4, 4), strides=(1, 1), padding="same"),
9 layers.LeakyReLU(alpha=0.2),
10 layers.GlobalMaxPooling2D(),
11 layers.Dense(1)
12
13 #generator's core
14 layers.Conv2DTranspose(64, (3, 3), strides=(2, 2), padding="same"),
15 layers.LeakyReLU(alpha=0.2),
16 layers.Conv2DTranspose(128, (3, 3), strides=(2, 2), padding="same"),
17 layers.LeakyReLU(alpha=0.2),
18 layers.Conv2DTranspose(256, (3, 3), strides=(2, 2), padding="same"),
19 layers.LeakyReLU(alpha=0.2),
20 layers.Conv2D(3, (4, 4), padding="same", activation="sigmoid")
```


Optimization

- ➡ From now on I will refer to **method two**, with which I obtained the best results.
- ➡ I've discovered the possibility to **customize** the `tf.keras.Model.fit` module and I used this approach for training the conditional GAN.
- ➡ For doing that, in a `cond_GAN` class, I **customized** three methods `metrics`, `compile` and `train_step`:

```
1 #metrics
2 #the mean value of the loss—fz obtained predicting the samples
3 self.gen_loss_tracker = keras.metrics.Mean(name="generator_loss")
4 self.disc_loss_tracker = keras.metrics.Mean(name="discriminator_loss")
5
6 def metrics(self):
7     return [self.gen_loss_tracker, self.disc_loss_tracker]
```

- ➡ Implementing an **exponential decay** of the learning rate [2] using `keras.optimizers.schedules.ExponentialDecay`;
- ➡ using a Binary Crossentropy as loss function.

```
1 #an exponential decay for the learning rate
2 initial_learning_rate = 0.0005
3
4 lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
5     initial_learning_rate, decay_steps=500, decay_rate=0.99, staircase=True)
6
7 #the decay's expression
8 def decayed_learning_rate(step):
9     return initial_learning_rate * decay_rate ^ (step / decay_steps)
10
11 cond_gan.compile(
12     d_optimizer=keras.optimizers.Adam(learning_rate=lr_schedule, beta_1=0.5),
13     g_optimizer=keras.optimizers.Adam(learning_rate=lr_schedule, beta_1=0.5),
14     loss_fn=keras.losses.BinaryCrossentropy(from_logits=True),
15 )
16
17 # Compile method is defined into the cGAN class
18 def compile(self, d_optimizer, g_optimizer, loss_fn):
19     super(ConditionalGAN, self).compile()
20     self.d_optimizer = d_optimizer      #Adam
21     self.g_optimizer = g_optimizer      #Adam
22     self.loss_fn = loss_fn              #BinaryCrossentropy
```

Learning rate decay

- ➡ My schedule performed a decay step each 50 optimization steps (after using 50 batches, so about every 2 epochs).



Figure 2: First 200 values of the learning rate used during the train in function of the epochs

Optimization - train step

The Algorithm's steps involve an **alternating training** of D and G . This procedure is done informing the optimizer about the correct label of the batches (0 real, 1 fake)⁴:

1. select a **real batch** of size n_{batch} ;
2. generate n_{batch} **latent points** (with both noise and label information inside) and use it for **generating fake images**;
3. use these **concatenated batches** for training only D .

```
1 def train_step(self, data):
2     #once selected the real batch, we prepare the fake one
3     random_latent_vectors = tf.random.normal(shape=(batch_size, self.latent_dim))
4     random_vector_labels = tf.concat([random_latent_vectors, one_hot_labels], axis=1)
5     generated_images = self.generator(random_vector_labels)
6
7     #let's concatenate images with their labels and create a big mixed batch
8     comb_images = tf.concat([fake_image_and_labels, real_image_and_labels], axis=0)
9     labels = tf.concat([tf.ones((batch_size, 1)), tf.zeros((batch_size, 1))], axis=0)
10
11    #training only the discriminator
12    with tf.GradientTape() as tape:
13        predictions = self.discriminator(comb_images)
14        d_loss = self.loss_fn(labels, predictions)
15        grads = tape.gradient(d_loss, self.discriminator.trainable_weights)
16        self.d_optimizer.apply_gradients(zip(grads, self.discriminator.trainable_weights))
```

⁴This is the opposite of the above, but this choice is conventional and due to the fact that `apply_gradients` performs a descent and not an ascent

4. generate other n_{batch} latent points and labels;
5. let G use them for generating a new **fake** sample;
6. let's train G on the cGAN predictions (here we *freeze* the discriminator)

```
1 #then we create another set of inputs for the generator
2 #associating them with 0 labels, because we want to deceive D
3     misleading_labels = tf.zeros((batch_size, 1))
4
5     with tf.GradientTape() as tape:
6         fake_images = self.generator(random_vector_labels)
7         fake_image_and_labels = tf.concat([fake_images, image_one_hot_labels], -1)
8         predictions = self.discriminator(fake_image_and_labels)
9         g_loss = self.loss_fn(misleading_labels, predictions)
10    grads = tape.gradient(g_loss, self.generator.trainable_weights)
11    self.g_optimizer.apply_gradients(zip(grads, self.generator.trainable_weights))
12
13    #updating the trackers state
14    self.gen_loss_tracker.update_state(g_loss)
15    self.disc_loss_tracker.update_state(d_loss)
16
17    return {
18        "g_loss": self.gen_loss_tracker.result(),
19        "d_loss": self.disc_loss_tracker.result(),
20    }
```

Results

Results

- ➡ I ran my code on Google Colab Pro, because such a problem requires a long lead time, even if well optimized. It provides **NVIDIA Tesla K80 GPUs** for a maximum single run of 12h;
- ➡ I also tried the kaggle GPUs, provided for a maximum of 37h of weekly usage. It provides **NVIDIA Tesla P100 PCI GPUs** for a maximum single run of 9h;

hyper-param	choice
n_{batch}	32
n_{epochs}	500
η_0	0.0005
γ_η	0.99
T_η	50
β_1	0.5
opt	Adam
J	Binary Crossentropy
z_{dim}	128

- ➡ the following result took $\sim 80'$ time.

➡ Fire line 1, Water line 2, Grass line 3.




Figure 3: Images before the training




Figure 4: Images after 10 epochs of training




Figure 5: Images after 100 epochs of training




Figure 6: Images after 200 epochs of training




Figure 7: Images after 500 epochs of training

🎬 cGAN learns to represent the three classes

➡ Thanks to [3] i've learnt that if we feed G with a series of *intermediate state* labels, we are able to generate *hybrid-classes' pokemon!*

```
1 #fire to grass label's transformation
2 #in this case we have 10 interpolations
3 tf.Tensor(
4 [[1.          0.          0.          ]
5  [0.8888889   0.11111111 0.          ]
6  [0.7777778   0.22222222 0.          ]
7  [0.6666666   0.33333334 0.          ]
8  [0.5555556   0.44444445 0.          ]
9  [0.44444442  0.5555556   0.          ]
10 [0.3333333   0.6666667   0.          ]
11 [0.22222221  0.7777778   0.          ]
12 [0.1111111   0.8888889   0.          ]
13 [0.          1.          0.          ]], shape=(10, 3), dtype=float32)
```

🎬 Fire to Grass transformation

🎬 Fire to Water transformation

A difficult training

- ➡ To train GANs, and cGANs too, is a difficult task
- ➡ Pokemon are very heterogeneous, even those belonging to the same type;
- ➡ during my job I performed several attempts desiring to achieve better and better results;
- ➡ I report the history of the loss functions evaluated during the last training:

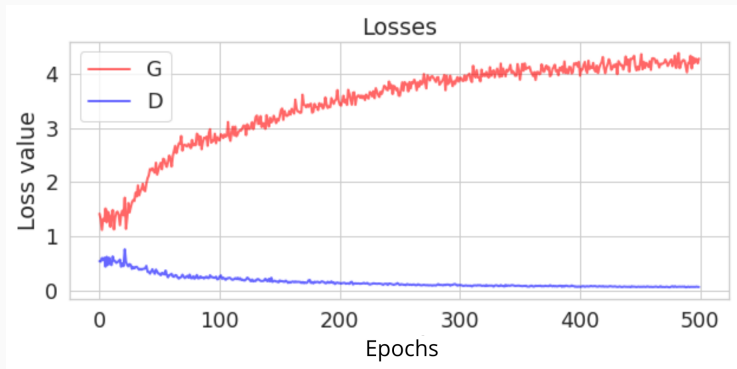


Figure 8: Value of loss function as a function of epochs

How to do better?

Data:

- searching for new data, maybe involving more than three classes and increasing the poor amount of images of my DS;
- building a dataset more balanced: in this case I had 351 fire pokes, 288 grass pokes and 425 water pokes.

Model:

- implement an hyper-optimization method, instead of doing that manually;
- test more complex models, in order to keep under control the training of G and D : for example, it's possible to use more than one encoder and more than one decoder;
- exploit verified models, such as the [NVIDIA Research Projects StyleGAN](#), using the power of the transfer learning;

Training:

- writing a more balanced train step method, for example training k times D and one time G , as said in [1], or training G with a batch of the same size of the one used with D .



Nolan J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, Yoshua Bengio, *Generative Adversarial Networks* (2014), [arXiv:1406.2661](#).



Mehdi Mirza, Simon Osindero, *Conditional Generative Adversarial Nets* (2014), [arXiv:1411.1784](#).



Sayak Paul, *conditional GAN* (2021), [official Keras' tutorial with MNIST](#).



Mikolaj Kolman, [Initial state of the pokemon's dataset](#)

Conv2DTranspose

